



Applied Deep Learning

Dr. Philippe Blaettchen
Bayes Business School (formerly Cass)

www.bayes.city.ac.uk

Learning objectives of today

Goals: Creating neural networks with Keras and TensorFlow

- Understanding the role that frameworks such as TensorFlow play in enabling the design, training, and use of arbitrarily complex neural networks
- Getting started on applying TensorFlow to create different types of feed-forward networks


How will we do this?

- We start with a quick recap of the functioning of neural networks
- We then introduce TensorFlow and Keras as programming frameworks for neural networks
- We create a number of neural networks, using TensorFlow's Sequential API




Recap

Components




... ..



Source: Liang

Hidden layers

$$“x” = a^{[l-1]} \quad z = a^{[l-1]}w + b \quad f(z)$$




- f is what we call an “activation function”
- There are many activation functions, and new ones are invented all the time
- Many of these functions do just fine, or slightly better than existing ones

Source: Liang


Typical activation functions: Rectified Linear Unit

Rectified Linear Unit (ReLU)



$$f(z) = \max\{0, z\}$$

“Derivative”




$$f'(z) = \begin{cases} 0, & \text{if } z < 0 \\ 1, & \text{if } z > 0 \end{cases}$$



Binary classification


- Input: $\mathbf{a}^{[L-1]}(i)$
- As usual, we make a linear transformation:
$$z^{[L]}(i) = \mathbf{a}^{[L-1]}(i) \mathbf{w}^{[L]} + b^{[L]}$$
- We then use the logistic sigmoid function
$$\hat{y}^{(i)} = f(z^{[L]}(i)) = \sigma(z^{[L]}(i)) = \frac{1}{1 + e^{-z^{[L]}(i)}}$$
- We can interpret the output as the probability of $y^{(i)} = 1$



Source: Liang

Multi-class classification

- We again make a linear transformation with a matrix of weights: $\mathbf{z}^{[L](i)} = \mathbf{a}^{[L-1](i)} \mathbf{W}^{[L]} + \mathbf{b}^{[L]}$
- Note that $\mathbf{z}^{[L](i)} = (z_1^{[L](i)} \quad z_2^{[L](i)} \quad \dots \quad z_K^{[L](i)})$
- We then use the softmax function on each of the outputs:
$$\hat{y}_k^{(i)} = f(\mathbf{z}^{[L](i)}) = \frac{e^{-z_k^{[L](i)}}}{\sum_{k=1}^K e^{-z_k^{[L](i)}}}$$
- This implies that $\hat{y}_k^{(i)} \in (0,1)$ and $\sum_{k=1}^K \hat{y}_k^{(i)} = 1$
- Hence, we can interpret $\hat{y}_k^{(i)}$ as the probability that $y^{(i)} = k$ (“belongs to class k ”)



Source: Liang

Generalizing our optimization algorithm

1. Decide a “learning rate” α
2. Start with some parameters θ
3. For a certain number of iterations
 - Compute $J(\theta)$
 - Let $\theta := \theta - \alpha \nabla_{\theta} J(\theta)$

(initialization)


(forward propagation)

(back-propagation)




BAYES
BUSINESS SCHOOL
CITY, UNIVERSITY OF LONDON

Step 1: Forward propagation through the computational graph



Step 2: Back-propagation through the computational graph





Deep learning programming frameworks

Deep learning in Python: the two main rivals

TensorFlow

- Developed by Google
- Python API based on C engine
- Largest community
- Parallelization of models and data
- Visualization (TensorBoard)
- Relatively intuitive when familiar with numpy
- Runs on most systems (e.g., TensorFlow Lite)
- Broader framework (e.g., reinforcement learning)
- Can be slower than other frameworks

Pytorch

- Developed by Facebook
- Python API based on LUA engine (Torch)
- Many pre-trained models
- Highly modular
- Easy to run on GPUs
- Runs on most systems (e.g., PyTorch Mobile)
- Lacking documentation


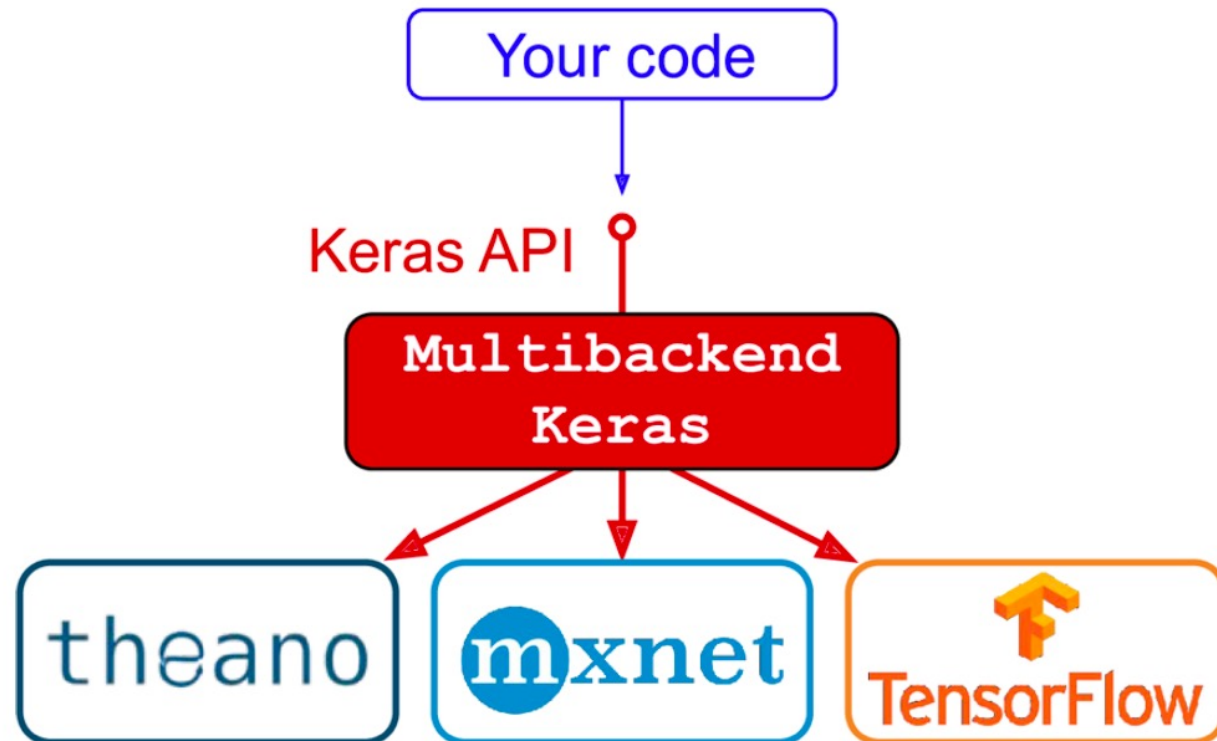


Wasn't there also this thing called Keras?

- High-level deep learning API to build, train, evaluate, execute neural networks
 - Documentation: <https://keras.io>
- Requires computation backend:
 - TensorFlow
 - Microsoft Cognitive Toolkit
 - Theano
 - Apache MXNET
 - Core ML (Apple)
 - JavaScript & TypeScript (for web browsers)
 - PlaidML
 - ...
- TensorFlow now comes with its own Keras implementation `tf.keras` (with some added features)
 - E.g., use Data API from TF



Keras in general and Keras using TensorFlow



Source: Geron



BAYES
BUSINESS SCHOOL
CITY, UNIVERSITY OF LONDON

TensorFlow in practice



BAYES
BUSINESS SCHOOL
CITY, UNIVERSITY OF LONDON



See you next week!

Sources

- DeepLearning.AI, n.d.: [deeplearning.ai](https://www.deeplearning.ai)
- Geron, 2019, Hands-on Machine Learning with Scikit-Learn, Keras, and TensorFlow
- Goodfellow, Bengio, Courville, 2016, The Deep Learning Book:
<http://www.deeplearningbook.org>
- Liang, 2016, Introduction to Deep Learning:
<https://www.cs.princeton.edu/courses/archive/spring16/cos495/>

